MPC for Group Reconstruction Circuits

Lúcás Críostóir Meier

June 21, 2022

Abstract

In this paper, we present a thing.

1 Introduction

Threshold Cryptography is important.

There have been protocols for Schnorr signatures, Threshold Encryption.

What makes these efficient is...

In this work we generalize this functionalities by...

We do an MPC protocol with commitments, and have a proof...

The essence of our protocol is...

Organization of the rest of this paper...

2 Background

Throughout this paper, we let $\mathbb{G}$ denote a group of prime order q, with generators G and H. Let $\mathbb{F}_q$ denote the scalar field associated with this group, and let $\mathbb{Z}/(q)$ denote the additive group of elements in this field. We also define $[n] := [1, \ldots, n]$.

We make heavy use of group homomorphisms throughout this paper. We let
$$\varphi(P_1, \ldots, P_m) : \mathbb{A} \to \mathbb{B}$$
denote a homomorphism from $\mathbb{A}$ to $\mathbb{B}$, parameterized by some public values $P_1, \ldots, P_m$. Commonly $\mathbb{A}$ will be a product of several groups $\mathbb{G}_1, \ldots, \mathbb{G}_n$, in which case we'd write:
$$\varphi(P_1, \ldots, P_m)(x_1, \ldots, x_n)$$
to denote the application of φ to an element $(x_1, \ldots, x_n)$ of the product group. The public values P_i are often left implicit.

We often write products $(x_1, \ldots, x_n)$ as a single vector $\mathbf{x} \in \mathbb{A}^n$. Operations between these vectors are done element-wise, so we write $\mathbf{x} + \mathbf{y}$ for $(x_1 + y_1, \ldots, x_n + y_n)$, as well as $\mathbf{x} \cdot G$ for $(x_1 \cdot G, \ldots, x_n \cdot G)$.

2.1 Pedersen Commitments

A key component of our scheme are Pedersen commitments [Ped92]. In their basic form, they allow one to commit to a value in $x \in \mathbb{Z}/(q)$. This is done by sampling a random $\alpha \xleftarrow{R} \mathbb{Z}/(q)$, and forming the commitment:

$$\text{Com}(x, \alpha) := x \cdot G + \alpha \cdot H$$

where H is a generator of $\mathbb{G}$, independent from G.

This scheme is *perfectly hiding*, because $\alpha \cdot H$ acts like a random element of $\mathbb{G}$, and completely masks $x \cdot G$.

On the other hand, this scheme is only *computationally* binding. This is because the discrete logarithm H with respect to G must be kept hidden. If the discrete logarithm of H is known, then it becomes possible to *equivocate*, by finding two different inputs (x, α) and (x', α') with equal commitments, i.e. $\text{Com}(x, \alpha) = \text{Com}(x', \alpha')$.

In fact, we can more precisely characterize this property: knowing the discrete logarithm of H is *necessary* in order to be able to equivocate.

Claim 2.1. Given two inputs $(x, \alpha) \neq (x', \alpha')$ such that $\text{Com}(x, \alpha) = \text{Com}(x', \alpha')$, it's possible to efficiently compute the discrete logarithm of H.

The proof is just a matter of algebra:

$$x \cdot G + \alpha \cdot H = x' \cdot G + \alpha' \cdot H$$
$$(x - x') \cdot G = (\alpha' - \alpha) \cdot H$$
$$\frac{(x - x')}{(\alpha' - \alpha)} \cdot G = H$$

Thus $(x - x')/(\alpha' - \alpha)$ is our discrete logarithm.

∎

2.1.1 Vector Pedersen Commitments

It's useful to generalize this scheme to the case of a vector of scalars $\mathbf{x} \in \mathbb{Z}/(q)^n$. The randomness becomes a vector of the same size, sampled as $\boldsymbol{\alpha} \xleftarrow{R} \mathbb{Z}/(q)^n$. We then define an analogous commitment scheme by doing scalar multiplication element-wise:

$$\text{Com}(\mathbf{x}, \boldsymbol{\alpha}) := \mathbf{x} \cdot G + \boldsymbol{\alpha} \cdot H$$

This generalization is naturally also perfectly hiding, and satisfies an analogous property with regards to equivocation:

Claim 2.2. Given two inputs $(\mathbf{x}, \boldsymbol{\alpha}) \neq (\mathbf{x}', \boldsymbol{\alpha}')$ such that $\mathrm{Com}(\mathbf{x}, \boldsymbol{\alpha}) = \mathrm{Com}(\mathbf{x}', \boldsymbol{\alpha}')$, it's possible to efficiently compute the discrete logarithm of H.

Since the two inputs are different, there must exist an index i such that $\mathbf{x}_i \neq \mathbf{x}'_i$ or $\boldsymbol{\alpha}_i \neq \boldsymbol{\alpha}'_i$. From here we apply Claim 2.1 with $(\mathbf{x}_i, \boldsymbol{\alpha}_i)$ and $(\mathbf{x}'_i, \boldsymbol{\alpha}'_i)$.

∎

2.1.2   On the Trusted Setup

In theory, Pedersen commitments require a trusted setup, to generate the group elements $G, H \in \mathbb{G}$. In practice, we argue that this trusted setup isn't a concern. This is the generator G is usually part of the specification for the group being used, and because there exist efficient methods to hash into elliptic curves [Ica09]. This reduces the problem of generating H to that of finding a credibly "unbiased" choice of seed to hash. This can be done in many ways.

One way would be to hash a canonical representation of G as bytes in order to produce H. Presumably, the generator G was not chosen in such a way as to produce an H with a known discrete logarithm, with that specific method of hashing into the group.

Another method would be to use a public source of randomness, such as public newspapers, lotteries [BDF+15], or Cryptographic protocols designed to provide such a service [FIP11].

Now, in practice while we can avoid a trusted setup with these methods, for our security analysis we do actually make use of this setup. Essentially, we prove that the security of our protocol reduces to the hardness of the discrete logarithm problem in $\mathbb{G}$, and to do this we need to be able to use an instance of the problem as a setup for the participants in our simulation.

2.2   Maurer's φ-Proof

In [Mau09], Maurer generalized Schnorr's sigma protocol for knowledge of the discrete logarithm [Sch90] to a much larger class of relations. In particular, Maurer provided a sigma protocol for proving knowledge of the pre-image of a group homomorphism φ. We denote this protocol as a "φ-proof", and recapitulate the scheme here.

Given a homomorphism $\varphi : \mathbb{A} \to \mathbb{B}$, and a public value $X \in \mathbb{B}$, the prover wants to demonstrate knowledge of a private value $x \in \mathbb{A}$ such that $\varphi(x) = X$. The prover does this by means of Protocol 2.1:

Protocol 2.1: φ-**Proof**

Prover	Verifier
knows $x \in \mathbb{A}$	public $X \in \mathbb{B}$

$$k \xleftarrow{R} \mathbb{A}$$
$$K \leftarrow \varphi(k)$$

$$\xrightarrow{\quad K \quad}$$

$$c \xleftarrow{R} \mathbb{Z}/(p)$$

$$\xleftarrow{\quad c \quad}$$

$$s \leftarrow k + c \cdot x$$

$$\xrightarrow{\quad s \quad}$$

$$\varphi(s) \stackrel{?}{=} K + c \cdot X$$

Here, p is chosen such that $\forall B \in \mathbb{B}.\ p \cdot B = 0$. In practice, we'll set $p = q$, which will work perfectly for the groups we use, which are all products of $\mathbb{G}$ or $\mathbb{Z}/(q)$.

Claim 2.3. Protocol 2.1 is a valid sigma protocol.

Completeness follows directly from the fact that φ is a homomorphism.

For the HVZK property, the simulator $\mathcal{S}(X, c)$ works by generating a random $s \xleftarrow{R} \mathbb{A}$, and then setting $K := \varphi(S) - c \cdot X$.

Finally, we prove 2-extractability. Given two verifying transcripts (K, c, s) and (K, c', s') sharing the first message, we extract a value $\hat{x}$ satisfying $\varphi(\hat{x}) = X$ as follows:

$$\varphi(s) - c \cdot X = K = \varphi(s') - c' \cdot X$$
$$\varphi(s) - \varphi(s') = c \cdot X - c' \cdot X$$
$$\frac{1}{c - c'} \cdot \varphi(s - s') = X$$
$$\varphi\left(\frac{s - s'}{c - c'}\right) = X$$

Thus, defining $\hat{x} := (s - s')/(c - c')$, we successfully extract a valid pre-image.

We conclude that the protocol is a valid sigma protocol.

■

Maurer's protocol can also work even in the case where the order of the groups are not known, but this makes the challenge generation a bit more complicated, and we don't need this functionality in this work.

2.3   Ideal Functionalities for Sigma Protocols

Functionality 2.1: Zero-Knowledge Functionality $\mathcal{F}(\texttt{ZK}, \varphi)$

A functionality $\mathcal{F}$ for parties $P_1, \ldots, P_n$.

On input $(\texttt{prove}, \texttt{sid}, x)$ from P_i:
$\mathcal{F}$ checks that $\texttt{sid}$ has not been used by P_i before.
$\mathcal{F}$ generates a new token π, and sets $x_\pi \leftarrow x$.
$\mathcal{F}$ replies with $(\texttt{proof}, \pi)$.

On input $(\texttt{verify}, X, \pi)$:
$\mathcal{F}$ replies with $(\texttt{verify-result}, \varphi(x_\pi) \stackrel{?}{=} X)$.

2.4   Broadcast Functionalities

The second ideal functionality we need is a *broadcast functionality*. The purpose of this functionality is to allow a party to send a message to all other parties, while guaranteeing that the party doesn't cheat by sending different messages.

Functionality 2.2: Authenticated Broadcast Functionality $\mathcal{C}$

A functionality $\mathcal{C}$ for parties $P_1, \ldots, P_n$.

On receiving $(\texttt{broadcast-in}, \texttt{sid}, m)$ from P_i:
$\mathcal{C}$ checks that $\texttt{sid}$ has not been used by P_i before.
$\mathcal{C}$ sends $(\texttt{broadcast-out}, \texttt{pid}_i, \texttt{sid}, m)$ to every party P_j.

In fact, our functionality only needs to be usable a single time, so it's in reality a *one-time* broadcast functionality, which might be simpler to implement.

The key difference between this broadcast functionality and simply sending a message to all other parties is that the functionality prevents malicious

participants from sending different messages to different parties. We use this later in the protocol to broadcast commitments, and in this situation it's important that all parties agree on what the values of these commitments are.

This functionality can be securely implemented in the UC framework using the "echo-broadcast" protocol from [GL05], which we recapitulate here:

Protocol 2.2: Echo-Broadcast Protocol

Each party P_i has a broadcast input m_i.

Each P_i sends $(\texttt{sid}, m_i)$ to all other parties.
Upon receiving $(\texttt{sid}, \hat{m}^j)$ from P_j, P_i checks that it hasn't already received a message from P_j for this $\texttt{sid}$, and then sends $(\texttt{rebroadcast}, \texttt{sid}, j, \hat{m}^j)$ to all other parties.
Upon receiving $(\texttt{rebroadcast}, \texttt{sid}, j, \hat{m}^j_i)$ from all parties, P_i checks that $\hat{m}^j_1 = \hat{m}^j_2 = \ldots = \hat{m}^j_n$, and then uses $\hat{m}^j_1$ as the message sent by P_j for this session.

This protocol UC securely implements Functionality 2.2.

We note that instead of each party forwarding $\hat{m}^j$, it's also possible to send $H(\hat{m}^j)$, where H is a collision-resistant hash function. This might be advantageous for long messages.

3 Group Reconstruction Circuits

If we look at the various functionalities we saw in the introduction, they all share quite a few commonalities. They first convert a threshold sharing of the input into a linear sharing of the input. These secret shared inputs can be added together, using $[x]$ and $[y]$ to form $[x + y]$, or used to multiply a group element, forming $[xA]$. These operations can be done locally. A secret shared value $[x]$ can be reconstructed, to have each party learn x.

In essence, *group reconstruction circuits* are a vast generalization of this kind of functionality. The core of the idea is that the operations done inside of these functionalities are in fact all group homomorphisms, with respect to all of their inputs. Because of these, the functionality can be efficiently computed by each party locally, with the reconstruction steps done by having each party reveal their share of the secret. Furthermore, in

the malicious setting, we can use φ-proofs in order to efficiently demonstrate the correctness of our computations.

3.1   Formal Definition

Formally, a group reconstruction circuit (GRC) is a special kind of directed acyclic graph (DAG), similar to an arithmetic circuit. This graph is *typed*, in the sense that each node is associated with some group $\mathbb{A}$, which represent the kinds of values that node is supposed to have. This graph features the following nodes:

- An input node, with type $\mathbb{Z}/(q)$.

- A random input node, type $\mathbb{Z}/(q)$.

- A reconstruction node, which has another node as input, and inherits the type of that node.

- A φ node, which can have several inputs, and which represents a homomorphism $\varphi(P_1, \ldots, P_l) : \mathbb{G}_1 \times \ldots \times \mathbb{G}_m \to \mathbb{H}$. For each of the public parameters P_i, there should be an input connected to a reconstruction node. For each of the inputs in $\mathbb{G}_i$, there should be an input connected to a node of that type.

- An output node, which has a single non-output node as input, inheriting its type.

Each φ node has a *depth* associated with it, equal to the largest number of reconstruction nodes on a path from an input node to the φ node, plus 1. We can define the depth of a circuit, by taking the largest depth among its φ nodes.

This circuit can be given semantics in the form of a semi-honest MPC protocol. For each of the input nodes with value x, the parties have a linear secret sharing $[\![x]\!]$. For random nodes, the parties sample a sharing $[\![k]\!]$ by locally generating a random share $k_i \xleftarrow{R} \mathbb{Z}/(q)$. For reconstruction nodes, the parties go from a secret sharing $[\![y]\!]$ to a public value y by having each party reveal their share y_i. For φ nodes, the parties use the fact that $[\![\varphi(x^1, \ldots, x^m)]\!] = \varphi([\![x^1]\!], \ldots, [\![x]\!]^m)$, by locally computing their share of the result as $\varphi(x_i^1, \ldots, x_i^m)$. Finally, for output nodes, they use their local share of a value, if that value is secret, or the public value, if it has been reconstructed. phrasing?

As an example, let's consider the functionality for Schnorr signatures.

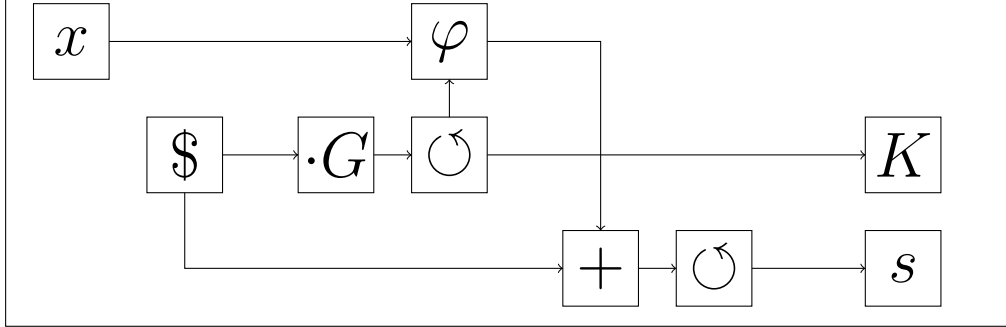7

Figure 3.1: **Schnorr Signature GRC**

Figure 3.1 presents the circuit for Schnorr signatures, with $\circlearrowleft$ denoting reconstruction nodes, \$ denoting random input nodes, and φ denoting the parameterized homomorphism:

$$\varphi(K)(x) := H(K, m) \cdot x$$

Both the φ and $+$ nodes have a depth of 2, and thus so does the circuit.

Using our semantics for circuits, we get a semi-honest protocol for threshold Schnorr signatures. The quorum turns their threshold secret sharing of the secret key x into a linear secret sharing $[\![x]\!]$. They then generate a shared nonce $[\![k]\!]$ by each sampling $k_i \overset{R}{\leftarrow} \mathbb{Z}/(q)$. They calculate a sharing $[\![K := k \cdot G]\!]$ by each computing $K_i := k_i \cdot G$. They then reveal these K_i to learn K. They can then compute $s_i := \varphi(K)(x_i)$ locally, creating a sharing $[\![s]\!]$, whose shares they then reveal, to learn s. They then use (K, s) as their signature.

3.2 Normalized Form

While the formal definition of GRCs above is complete, and closely matches the description of various functionalities, it's not particularly convenient to design a protocol for. We can vastly simplify the description of a GRC through the use of a *normalized form*, which is a much more compact representation of this circuit. This normalized form is also directly amenable to an implementation as an MPC protocol. This form is derived from a series of simple transformations on a GRC.

First, note that the random input nodes of the GRC don't depend on any other node. Because of this, we can move all of these input nodes to the start of the circuit, which has the semantics of generating all the randomness in the circuit at the start of the execution.

Second, instead of having several input nodes $x^1, \ldots, x^m$, all elements of $\mathbb{Z}/(q)$, we instead have a single input *vector* $\mathbf{x} \in \mathbb{Z}/(q)^m$. This vector can also be considered as a single element of the product group, with addition defined pointwise, as in Section 2. In fact, we can also include random elements as part of the input. The idea is that honest parties will generate this part of the input randomly for each execution. We treat this issue in more detail in Section 4.1.

Third, we can make each φ node depend on the entire input vector $\mathbf{x}$. In practice, this dependency can be sparse, in that φ will only make use of a small number of elements in the input. Nonetheless, φ is still a homomorphism with respect to the entire input. This is because a projection $\pi : \mathbb{Z}/(q)^m \to \mathbb{Z}/(q)^l$, which selects a subset of elements from an input vector, is a group homomorphism. Any homomorphism using only a subset of elements can be composed with π to make a homomorphism taking in the entire vector.

Fourth, we can coalesce the homomorphisms together, creating one homomorphism for each "layer" of the circuit. We do this by first organizing the φ nodes into layers, based on their depth. Each node with the same depth goes into the same layer. We then coalesce all of the homomorphisms in this layer into a single homomorphism φ. We can do this because the duplication map $a \mapsto (a, a)$ and the projection map $(a, \_) \mapsto a$ are both homomorphisms. We can sort all of the homomorphisms topologically, and then compose them sequentially, duplicating the input and projecting as necessary.

We can also remove reconstruction nodes. Because each layer only has a single homomorphism, we can consider the output of this homomorphism to neccessarily be reconstructed. We then make each homomorphism parameterized by all of the reconstructed outputs from each previous layer.

Finally, we can remove output nodes. By considering the output of every layer to be part of the output, we include all of the output nodes connected to reconstruction nodes. For the other outputs, they can be locally computed using the reconstructed outputs, as well as the shares of the input vector, so it's not necessary to include them in the circuit.

Combining all of these gives us the formal description of normalized form GRCs in Figure 3.2.

We can also give circuits in this form semantics, in the form of a semi-honest MPC protocol. The parties have a linear secret sharing $[\![\mathbf{x}]\!]$ of the input vector, some entries of which having been generated randomly for this execution. Then, for each layer $r \in 1, \ldots, d$, the parties locally compute $\mathbf{V}_r^i := \varphi_r(\mathbf{V}_1, \ldots, \mathbf{V}_{r-1})$, and then reveal these shares, allowing each party to compute $\mathbf{V}_r := \sum_i \mathbf{V}_r^i$. The values $\mathbf{V}_1, \ldots, \mathbf{V}_d$ make up the output of the protocol.

In Section 5, we provide examples of GRCs in normalized form for several functionalities, including Schnorr signatures.

4 MPC Protocol for GRCs

In this section, we describe an MPC protocol for computing a GRC on linearly shared inputs, with associated commitments. We also analyze the security of this protocol, proving that it is secure against an arbitrary number of malicious parties, and under concurrent composition, in the UC framework.

For inputs, one natural kind of commitment are Pedersen commitments. In many protocols, however, it's more natural to use plain commitments, where a scalar x is committed to with the value $x \cdot G$. This matches most threshold Schnorr schemes, in which the secret x is split into shares x_i, with the shares of the public key $X_i := x_i \cdot G$ being known for all parties. While we could subsume these commitments as a case of Pedersen commitments, with a blinding factor set to 0, explicitly considering these plain commitments yields a more efficient protocol.

We thus split our input vector into three sections: $\mathbf{x}, \mathbf{y}, \mathbf{k}$. Each of this is linearly split into shares. We have shares $\mathbf{x}^1, \ldots, \mathbf{x}^n$ for each party, such that $\mathbf{x} = \sum_i \mathbf{x}^i$, and similarly for $\mathbf{y}$ and $\mathbf{k}$. The shares of $\mathbf{x}$ have plain commitments $\mathbf{X}^i = \mathbf{x}_i \cdot G$ for each party. The shares of $\mathbf{y}$ have Pedersen commitments $\mathbf{Y}^i = \mathbf{y}_i \cdot G + \boldsymbol{\alpha}^i \cdot H$, with $\boldsymbol{\alpha}^i$ a vector of blinding factors held by each party. Finally, $\mathbf{k}$ is intended to be randomly generated for

each execution of the protocol. Honest parties will generate their share $\mathbf{k}^i$ by sampling a random vector. As long as at least one participant in the protocol is honest, then $\mathbf{k} := \sum_i \mathbf{k}^i$ will also be random.

4.1 Ideal Functionality

In this section, we describe an ideal functionality for our protocol, as Functionality 4.1. This functionality is parameterized by the circuit Φ, as well as the input commitments $\mathbf{X}^i$ and $\mathbf{Y}^i$, for each party $i \in [n]$. The functionality also uses a common reference string $(G, H) \in \mathbb{G}^2$, for Pedersen commitments.

Functionality 4.1: GRC functionality $\mathcal{F}(\texttt{GRC}, \Phi, \mathbf{X}^i, \mathbf{Y}^i)$

A functionality $\mathcal{F}$ for parties $P_1, \ldots, P_n$.

After receiving $(\texttt{input}, \texttt{sid}, \mathbf{x}^i, \mathbf{y}^i, \boldsymbol{\alpha}^i, \mathbf{k}^i)$ from every party P_i:
$\mathcal{F}$ checks, for every $i \in [n]$, that:

$$\mathbf{X}^i \stackrel{?}{=} \mathbf{x}^i \cdot G$$

$$\mathbf{Y}^i \stackrel{?}{=} \mathbf{y}^i \cdot G + \boldsymbol{\alpha}^i \cdot H$$

$\mathcal{F}$ computes, for each round $r \in [d]$:

$$\mathbf{V}_r^i := \varphi_r(\mathbf{V}_1, \ldots, \mathbf{V}_{r-1})(\mathbf{x}^i, \mathbf{y}^i, \mathbf{k}^i)$$

$$\mathbf{V}_r := \sum_j \mathbf{V}_r^i$$

$\mathcal{F}$ sends $(\texttt{output}, \texttt{sid}, \mathbf{V}_1^1, \ldots, \mathbf{V}_d^n)$ to every party P_i.

This functionality checks that the inputs each party provides match the public commitments, and then computes the output of the circuit in a straightforward manner. One slight difference is that instead of simply learning $\mathbf{V}_r$ for every round r, each party learns $\mathbf{V}_r^i$ for every party i. Naturally, we have $\mathbf{V}_r = \sum_i \mathbf{V}_r^i$, so this information can be derived by each party. The reason we allow the parties to also learn the individual shares being revealed is that our eventual protocol will also reveal this information, so we need to model the leakage in our functionality as well. Furthermore, this matches the semantics of group reconstruction circuits, where parties learn the individual shares of the group element they're reconstructing. For practical functionalities like Schnorr signatures or threshold encryption, learning these intermediate values is not a concern either.

4.2 Protocol

In this section, we provide a protocol implementing Functionality 4.1.

The basic idea is that for each round r, the parties locally compute $\mathbf{V}_r^i :=$ $\varphi_r(\mathbf{x}^i, \mathbf{y}^i, \mathbf{k}^i)$, and then send these values to other parties, along with a proof that the value was computed correctly, and that the inputs used correspond to the public commitments.

For the random input $\mathbf{k}^i$, we need to guarantee that the same input vector is used throughout the protocol. We do this by creating a Pedersen commitment $\mathbf{K}^i$ to the random value, and having an initial round where each party broadcasts this commitment to the other parties. To prevent a party from sending different commitments, we use Functionality 2.2 to guarantee that the same commitment is sent to all parties.

We can easily prove that each step was computed correctly, and with the right inputs, by using the following homomorphism:

$$\psi_r(\mathbf{x}, \mathbf{y}, \boldsymbol{\alpha}, \mathbf{k}, \boldsymbol{\beta}) := (\varphi_r(\mathbf{x}, \mathbf{y}, \mathbf{k}),\ \mathbf{x} \cdot G,\ \mathbf{y} \cdot G + \boldsymbol{\alpha} \cdot H,\ \mathbf{k} \cdot G + \boldsymbol{\beta} \cdot H)$$

This homomorphism uses the same inputs to compute φ_r, but also reconstructs all of the commitments to the inputs. The φ-proofs seen in Section 2.2 provide an efficient sigma protocol to verify that a value $\mathbf{V}_r^i$ was computed using φ_r on the correct inputs. We then use Functionality 2.1 to turn these sigma protocols into ZK proof of knowledge functionalities.

Protocol 4.1 describes all of this more formally. Like the ideal functionality, the protocol is parameterized by the circuit Φ, the public commitments $\mathbf{X}^i, \mathbf{Y}^i$, and makes use of a common reference string $(G, H) \in \mathbb{G}^2$, for Pedersen commitments. The protocol takes $d + 1$ rounds, with d the depth of the circuit. As we mentioned in the previous section, the parties learn the intermediate values $\mathbf{V}_r^i$ as a consequence of the protocol's execution.

<div style="border: 1px solid black; padding: 10px;">

Protocol 4.1: MPC protocol for $\Phi, \mathbf{X}^i, \mathbf{Y}^i$

Each party P_i has inputs $\mathbf{x}^i$ and $\mathbf{y}^i$, committed to by $\mathbf{X}^i$ and $\mathbf{Y}^i$. They also have decommitments $\boldsymbol{\alpha}^i$ for $\mathbf{Y}^i$. Each party P_i also has a vector $\mathbf{k}^i$, which honest parties will have generated randomly.

Round 0
Each party P_i generates a random vector $\boldsymbol{\beta}^i$, and creates a commitment to $\mathbf{k}^i$ with:
$$\mathbf{K}^i := \mathbf{k}^i \cdot G + \boldsymbol{\beta}^i \cdot H$$

P_i sends $(\texttt{broadcast-in}, \texttt{sid}, \mathbf{K}^i)$ to the broadcast functionality $\mathcal{C}$.
P_i waits to receive $(\texttt{broadcast-out}, \texttt{sid}, \mathbf{K}^j)$ for each other party j.

Round r
Each party P_i computes $\mathbf{V}_r^i := \varphi_r(\mathbf{V}_1, \ldots, \mathbf{V}_{r-1})(\mathbf{x}^i, \mathbf{y}^i, \mathbf{k}^i)$.
Each party P_i sends $(\texttt{prove}, \texttt{sid}, (\mathbf{x}^i, \mathbf{y}^i, \boldsymbol{\alpha}^i, \mathbf{k}^i, \boldsymbol{\beta}^i))$ to $\mathcal{F}(\texttt{ZK}, \psi_r)$, receiving π_r^i in return.
Each party P_i sends $(\mathbf{V}_r^i, \pi_r^i)$ to every other party.

After receiving $(\mathbf{V}_r^j, \pi_r^j)$ from all other parties, P_i checks, for each j, that the proof is valid, by sending $(\texttt{verify}, (\mathbf{V}_r^j, \mathbf{X}^j, \mathbf{Y}^j, \mathbf{K}^j), \pi_r^j)$ to $\mathcal{F}(\texttt{ZK}, \psi_r)$, and aborting if the functionality returns 0.
Each party P_i then stores each $\mathbf{V}_r^j$ as part of its output, and computes $\mathbf{V}_r := \sum_j \mathbf{V}_r^j$.

</div>

4.3 Security Analysis

In this section, we prove that Protocol 4.1 implements Functionality 4.1 with UC security, even against an arbitrary number of malicious parties. More specifically, we work with the SUC model of security [CCL15]. We also work in the hybrid model, using Functionalities 2.1 and 2.2 for ZK proofs of knowledge, and authenticated broadcast, respectively. We also need a common reference string $(G, H) \in \mathbb{G}^2$, for Pedersen commitments. In our proof, the simulator provides this string.

Claim 4.1. Provided that the discrete logarithm is hard in $\mathbb{G}$, Protocol 4.1 securely implements Functionality 4.1, in the hybrid model of universally composable security, given a zk functionality $\mathcal{F}(\texttt{ZK}, \varphi)$ (for arbitrary φ), a broadcast functionality $\mathcal{C}$, as well as a common reference string $(G, H) \in \mathbb{G}^2$.

Proof Idea:
The basic idea is that we use an instance of the discrete logarithm problem

as our common reference string. This makes any break of the binding property of Pedersen commitments yield a solution to the discrete logarithm. Otherwise, we can perfectly simulate an execution of the protocol using the ideal functionality, without rewinding the adversary, which lets us conclude that our protocol is UC secure, using the result in [KLR09].

One technical detail is that we need inputs from every party to give to the ideal functionality. Because of this, our simulator first runs the simulated protocol until the adversary provides the input for their parties in the first ZK proof. We can then use these values for the parties the simulator controls in the ideal functionality.

The output of this functionality gives us all of the values we need in the rest of the simulation, except for the commitments $\mathbf{K}^i$ to the random inputs of the honest parties. For these, we use the fact that Pedersen commitments are perfectly hiding, and simply generate them at random.

Proof:

We prove this by constructing a simulator $\mathcal{S}$ which uses the ideal functionality $\mathcal{F}(\texttt{GRC})$ to perfectly simulate an execution of the hybrid protocol against an adversary $\mathcal{A}$.

We also work in the common reference string model, where the simulator $\mathcal{S}$ chooses the bases (G, H) for the Pedersen commitments.

We use this simulator $\mathcal{S}$ to construct an adversary against the discrete logarithm game.

Let $\mathcal{M} \subseteq \mathcal{P}$ be the set of malicious parties, and $\mathcal{H} \subseteq \mathcal{P}$ be the set of honest parties. Naturally, we have $\mathcal{H} \cup \mathcal{M} = \mathcal{P}$ and $\mathcal{H} \cap \mathcal{M} = \emptyset$.

As an adversary against the discrete logarithm game, $\mathcal{S}$ receives (G, H) as an instance of the discrete logarithm problem.

The simulator then proceeds as follows:
$\mathcal{S}$ starts by setting (G, H) as the common reference string.

Round 0:
For each $i \in \mathcal{H}$, $\mathcal{S}$ samples $\mathbf{K}^i \xleftarrow{R} \mathbb{G}$.
For each $i \in \mathcal{M}$, $\mathcal{S}$ waits to receive $(\texttt{broadcast-in}, \texttt{sid}, \mathbf{K}^i)$.
$\mathcal{S}$ then sends $(\texttt{broadcast-out}, \texttt{pid}_i, \texttt{sid}, \mathbf{K}^i)$, to all parties, for every $i \in \mathcal{P}$, emulating $\mathcal{C}$.

Interim:
$\mathcal{S}$ waits to receive $(\texttt{prove}, \texttt{sid}, (\mathbf{x}^i, \mathbf{y}^i, \boldsymbol{\alpha}^i, \mathbf{k}^i, \boldsymbol{\beta}^i))$ for each malicious $i \in \mathcal{M}$, playing the role of $\mathcal{F}(\texttt{ZK}, \psi_1)$.

$\mathcal{S}$ checks, for each i, that:

$$\mathbf{X}^i \stackrel{?}{=} \mathbf{x}^i \cdot G$$

$$\mathbf{Y}^i \stackrel{?}{=} \mathbf{y}^i \cdot G + \boldsymbol{\alpha}^i \cdot H$$

$$\mathbf{K}^i \stackrel{?}{=} \mathbf{k}^i \cdot G + \boldsymbol{\beta}^i \cdot H$$

otherwise, $\mathcal{S}$ sets $\texttt{bad-values}_1^i \leftarrow 1$.
$\mathcal{S}$ records the values $\mathbf{x}^i, \mathbf{y}^i, \boldsymbol{\alpha}^i, \mathbf{k}^i, \boldsymbol{\alpha}^i$, for $i \in \mathcal{M}$.

Now, in the real execution against $\mathcal{F}(\texttt{GRC})$, with real honest parties P_i, for each $i \in \mathcal{M}$, the parties $\mathcal{S}$ controls, $\mathcal{S}$ sends $(\texttt{input}, \texttt{sid}, \mathbf{x}^i, \mathbf{y}^i, \boldsymbol{\alpha}^i, \mathbf{k}^i)$ to $\mathcal{F}(\texttt{GRC})$.
$\mathcal{S}$ receives $(\texttt{output}, \texttt{sid}, \mathbf{V}_1^1, \ldots, \mathbf{V}_d^n)$ in return, and records these values.

Round r:
For each round $r \in [d]$, $\mathcal{S}$ proceeds as follows:
$\mathcal{S}$ generates a new π_r^j for each $j \in \mathcal{H}$, and sends $(\mathbf{V}_r^j, \pi_r^j)$ to every malicious P_i, with $i \in \mathcal{M}$.

Unless $r = 1$, $\mathcal{S}$ waits to receive $(\texttt{prove}, \texttt{sid}, \hat{\mathbf{x}}^i, \hat{\mathbf{y}}^i, \hat{\boldsymbol{\alpha}}^i, \hat{\mathbf{k}}^i, \hat{\boldsymbol{\beta}}^i)$ from each malicious P_i, for $i \in \mathcal{M}$, playing the role of $\mathcal{F}(\texttt{ZK}, \psi_r)$.
$\mathcal{S}$ then checks, for each i, that:

$$\mathbf{X}^i \stackrel{?}{=} \hat{\mathbf{x}}^i \cdot G$$

$$\mathbf{Y}^i \stackrel{?}{=} \hat{\mathbf{y}}^i \cdot G + \hat{\boldsymbol{\alpha}}^i \cdot H$$

$$\mathbf{K}^i \stackrel{?}{=} \hat{\mathbf{k}}^i \cdot G + \hat{\boldsymbol{\beta}}^i \cdot H$$

and sets $\texttt{bad-values}_r^i \leftarrow 1$ otherwise.
The first check implies that $\hat{\mathbf{x}}^i = \mathbf{x}^i$. If it holds that $\hat{\mathbf{y}}^i \neq \mathbf{y}^i$ or $\hat{\mathbf{k}}^i \neq \mathbf{k}^i$, then $\mathcal{S}$ has found a value h such that $h \cdot G = H$, as per Claim 2.2, and $\mathcal{S}$ aborts, returning h.

(Including when $r = 1$) $\mathcal{S}$ generates a new π_r^i, and returns $(\texttt{proof}, \pi_r^i)$, playing the role of $\mathcal{F}(\texttt{ZK}, \psi_r)$.

Concurrently, $\mathcal{S}$ plays the role of $\mathcal{F}(\texttt{ZK}, \psi_r)$, responding to $(\texttt{verify}, (\hat{\mathbf{V}}_r^j, \hat{\mathbf{X}}^j, \hat{\mathbf{Y}}^j, \hat{\mathbf{K}}^j), \pi)$ queries. $\mathcal{S}$ checks that there exists some $i \in \mathcal{P}$ such that $\pi_r^i = \pi$. $\mathcal{S}$ then returns:

$$\hat{\mathbf{V}}_r^j \stackrel{?}{=} \mathbf{V}_r^j \wedge \hat{\mathbf{X}}^j \stackrel{?}{=} \mathbf{X}^j \wedge \hat{\mathbf{Y}}^j \stackrel{?}{=} \mathbf{Y}^j \wedge \hat{\mathbf{K}}^j \stackrel{?}{=} \mathbf{K}^j \wedge \texttt{bad-values}_r^i \neq 1$$

$\mathcal{S}$ then waits to receive $(\hat{\mathbf{V}}^i, \hat{\pi}_r^i)$ for every malicious party P_i, with $i \in \mathcal{M}$.
$\mathcal{S}$ then checks if the query $(\texttt{verify}, (\hat{\mathbf{V}}_r^i, \mathbf{X}^i, \mathbf{Y}^i, \mathbf{K}^i), \hat{\pi}_r^i)$ would yield 1,

15

according to the logic in the section above. (If $\hat{\pi}_r^i$ doesn't match anything, the check is considered to fail). If this check fails, then $\mathcal{S}$ simulates every honest P_j aborting, with $j \in \mathcal{H}$, to abort, as if they'd seen an invalid proof themselves.

This concludes the simulation.

If $\mathcal{S}$ aborts with a value h, then they've successfully solved an instance of the discrete logarithm problem. Under our assumption that this problem is hard, this happens with negligeable probability.

We argue that if $\mathcal{S}$ does not abort in this way, then the simulation is perfect. For the first round, because pedersen commitments are perfectly hiding, sampling a random $\mathbf{K}^i$ has an identical distribution as an honest party generating a pedersen commitment. For the rest of the protocol, all of our checks are equivalent to those made by honest parties. This is because the $\mathbf{V}_i^j$ values are necessarily computed correctly, and use the inputs provided by the parties the adversary $\mathcal{A}$ controls.

Because our simulator $\mathcal{S}$ is perfect, and doesn't rewind the adversary $\mathcal{A}$, we conclude, using [KLR09], that our protocol satisfies universally composable security, in the hybrid model.

∎

4.3.1 Identifiable Abort

We note that our protocol can be considered to have identifiable aborts, provided that the implementation of the broadcast functionality also has this property. After round 0, if a party causes an abort by providing an invalid proof, then that abort can be detected, since there will be a signed message containing this invalid proof, which can be used as evidence of cheating.

4.4 Practical Considerations

In this section we describe a few additional details which might be useful for concrete implementations of the protocol.

4.4.1 Sparse Proofs

For simplicity, we described each homomorphism φ_i in the circuit as taking in the entire input vector $\mathbf{x}$. These homomorphisms may be *sparse*, in the sense that they only use a few of the elements in the input vector. It should be noted that in this case, the φ-proofs can be done as if only these select elements existed, which gives a slightly smaller and faster proof.

4.4.2 Removing Random Commitments

Some functionalities may not use any random input **k**. In this case, the first round can be omitted, since there's no need to commit to an empty vector. This actually improves the complexity of threshold encryption, as we'll see in Section cite.

4.4.3 Parallel Echos

In the description of our protocol, we use one round for broadcasting the commitments to the random inputs, using our broadcast functionality. If we naively replace this functionality with the echo-broadcast in Protocol 2.2, we increase the round complexity, since that protocol requires 2 rounds. However, the second round of the echo-broadcast protocol is only need to check the consistency of the broadcasted input; the parties already know the input after the first round. Because of this, we can run the re-broadcast round *in parallel* with the rest of the protocol, and thus incur no additional cost to round complexity.

4.4.4 Parallel `sid`

While there are different ways of instantiating the ZK proofs of knowledge for Functionality 2.1, some of them heavily rely on the use of a unique session identifier `sid`. Because this particular identifier isn't needed in the first round, we can run a session identifier agreement protocol, such as [BLR04], in parallel with the first round. This improves the round complexity of the concrete protocol as compared to sequential composition.

4.5 Complexity

The round complexity of our protocol is near-optimal, with only one additional round to provide commitments to random inputs, which can be omitted if the circuit doesn't use them.

We also argue that the concrete complexity of our protocol compares favorably with more specialized protocols. As noted before, the basic approach of locally computing the homomorphisms matches the basic strategy of specialized protocols. The overhead in our case comes from the additional φ-proofs we need to include, as well as computing additional commitments for the random inputs. Thankfully, these proofs are quite efficient. The cost of each proof is dominated by that of computing φ, along with 5 scalar multiplications, to prove that the right inputs were used. While more expensive than more specialized protocols, this is nonetheless a manageable amount of overhead for a generic protocol.

5 Applications

In this section, we describe several applications of our protocol, by providing concrete group reconstruction circuits for various functionalities. We also compare the complexity of our generic protocol, instantiated on these circuits, with specialized protocols computing the same functionality.

One example application for DKG, Schnorr, Threshold Encryption, Polynomial Commitments

5.1 Distributed Key Generation

5.2 Threshold Schnorr Signatures

We've described the Schnorr signature functionality as a GRC before, but not yet in normalized form. In this form, our input consists of the private signature key x, as well as the nonce k. We then have two homomorphisms:

$$\varphi_1(x, k) := k \cdot G$$
$$\varphi_2(K)(x, k) := k + H(K, m) \cdot x$$

We can then instantiate our protocol with this circuit to get a threshold signature scheme, assuming that some other distributed key generation scheme is used. Given a threshold sharing of the private key x, we assume that a quorum of parties can generate a linear sharing of x as $x_1 + \ldots + x_n$, along with public key shares $X^i := x_i \cdot G$. This assumption matches the standard polynomial schemes used for threshold secret sharing, where Lagrange interpolation is used to derive linear shares. Using these linear shares, the parties then run our protocol to compute a signature.

The round complexity of this signing protocol matches that 3 round scheme of Lindell [Lin22]. The only additional security assumption we make is that hardness of the discrete logarithm in $\mathbb{G}$, which makes our protocol relatively conservative, as Lindell's scheme purports to be. The concrete complexity of our protocol is a bit higher, because of the extra proofs we need to compute, and our use of Pedersen commitments.

Our protocol compares disfavorably against 2 round signing protocols such as FROST [KG20], or MuSig2 [NRS21], which also have better concrete complexity.

5.3 Threshold Encryption

5.4 Polynomial Commitments and Openings

6 Limitations and Further Work

Aggregating proofs

Exploiting circuit structure for lighter proofs

Threshold Interpretations

7 Conclusion

Summarize stuff

References

[BDF+15] Thomas Baignères, Cécile Delerablée, Matthieu Finiasz, Louis Goubin, Tancrède Lepoint, and Matthieu Rivain. Trap Me If You Can – Million Dollar Curve. 2015.

[BLR04]   Boaz Barak, Yehuda Lindell, and Tal Rabin. Protocol Initialization for the Framework of Universal Composability. 2004.

[CCL15]   Ran Canetti, Asaf Cohen, and Yehuda Lindell. A Simpler Variant of Universally Composable Security for Standard Multiparty Computation. In *CRYPTO 2015*, volume 9216, pages 3–22. Springer, Berlin, Heidelberg, 2015.

[FIP11]   Michael J. Fischer, Michaela Iorga, and René Peralta. A public randomness service. pages 434–438, July 2011.

[GL05]   Shafi Goldwasser and Yehuda Lindell. Secure Multi-Party Computation without Agreement. *Journal of Cryptology*, 18(3):247–287, July 2005.

[Ica09]   Thomas Icart. How to Hash into Elliptic Curves. In *CRYPTO 2009*, volume 5677 of *LNCS*, pages 303–316. Springer, Berlin, Heidelberg, 2009.

[KG20]   Chelsea Komlo and Ian Goldberg. FROST: Flexible Round-Optimized Schnorr Threshold Signatures. *Cryptology ePrint Archive*, 2020.

[KLR09]   Eyal Kushilevitz, Yehuda Lindell, and Tal Rabin. Information-Theoretically Secure Protocols and Security Under Composition. 2009.

[Lin22] Yehuda Lindell. Simple Three-Round Multiparty Schnorr Signing with Full Simulatability. *Cryptology ePrint Archive*, 2022.

[Mau09] Ueli Maurer. Unifying Zero-Knowledge Proofs of Knowledge. In *AFRICACRYPT 2009*, volume 5580 of *LNCS*, pages 272–286. Springer, Berlin, Heidelberg, 2009.

[NRS21] Jonas Nick, Tim Ruffing, and Yannick Seurin. MuSig2: Simple Two-Round Schnorr Multi-signatures. In *CRYPTO 2021*, volume 12825 of *LNCS*, pages 189–221. Springer, 2021.

[Ped92] Torben Pryds Pedersen. Non-Interactive and Information-Theoretic Secure Verifiable Secret Sharing. In *CRYPTO 91*, volume 576 of *LNCS*, pages 129–140. Springer, Berlin, Heidelberg, 1992.

[Sch90] C. P. Schnorr. Efficient Identification and Signatures for Smart Cards. In *CRYPTO 1989*, volume 435 of *LNCS*, pages 239–252, New York, NY, 1990. Springer.